

Self driving toy car

Project for 3D Computer Vision lecture, summer term 2020

Alexander Barth

*M. Sc. student computer engineering
Heidelberg University
dl248@stud.uni-heidelberg.de*

Grewan Hassan

*M. Sc. student physics
Heidelberg University
g.hassan@stud.uni-heidelberg.de*

Denis Münch

*M. Sc. student applied computer science
Heidelberg University
denis.muensch@stud.uni-heidelberg.de*

Royden Wagner

*M. Sc. student computer engineering
Heidelberg University
royden-wagner@outlook.com*

Abstract—This report describes the self driving car toy project done in the 3D Computer Vision lecture at Heidelberg University. The goal is to train neural networks so that the given car can drive autonomously on a track.

Index Terms—autonomous driving, computer vision, image processing, neural networks

I. GETTING STARTED

For getting started an operating system needs to be flashed onto the Raspberry Pi 3 B+ which is mounted into the car. Through the Raspberry Pi Imager the Pi OS 32-bit in release 2020-05-27 was flashed onto the SD card. The OS is a port of Debian with the Raspberry Pi Desktop and comes with an integrated configurator to enable SSH, VNC, the camera, SPI and I2C. During our project several different approaches are used. Roughly divided into an approach with and without machine learning. The first approach is classical image processing without machine learning by using OpenCV.

II. IMAGE PROCESSING WITHOUT ML

The goal of this approach is to let the car drive autonomously on the track without using machine learning. Therefore the approach can be roughly divided into two subtasks, the image processing with OpenCV and the driving. OpenCV is an open library for computer vision and machine learning of which the first part is used.

A. Driving

Once again the driving task can be divided into the following two subtasks, steering and throttle. The Raspberry is mounted on the vehicle together with a servo and an Arduino PCA9685 controller. The donkeycar application is steering through the PCA9685. This can be observed while following the donkeycar user guide in the step of calibrating the car. Calibrating the donkeycar takes the plugged in channel and the bus, which are the same required options using the *Adafruit_9685* library. Every time the calibrate function is executed a new donkeycar instance is initiated. Therefore our driving function writes directly to the controller without the donkeycar instance. Nevertheless the donkeycar calibrate

TABLE I: PWM values

mode	value
left	460
center	380
right	290
forwards	500
stop	370
backwards	220

function is helpful for getting the PWM values at which the steering is full left, straight and full right. Using the other channel the same procedure can be used to get the PWM values for full throttle, zero and full throttle backwards. Table I shows the different PWM values observed while using the calibrate function.

By using the boundary values each with the null value a respective linear equation can be set up.

$$\text{Left} \quad \text{pulse_length} = -4.4 * \text{angle} + 380$$

$$\text{Right} \quad \text{pulse_length} = -3.6 * \text{angle} + 380$$

$$\text{Forwards} \quad \text{pulse_length} = 1.4 * \text{acceleration} + 360$$

$$\text{Backwards} \quad \text{pulse_length} = 0.6 * \text{acceleration} + 360$$

These resulting equations are used in *driving_functions.py*. It can be used for steering and controlling the motor directly and is later used inside the python script for autonomous driving. The input for the steering function are steering angles reaching from -25 to 25 which are converted into degrees and into PWM signals. Using steering angles is easier for the human because they do not need to be abstracted. Similar way is used for the motor control. Input values to the motor control function are percentages of the acceleration reaching from -100 to 100 which are again converted into PWM signals.

B. Lane Detection

Second subtask for driving autonomously is to detect the track lanes. The result is shown in Figure 1. Figure 1d shows the input frame which is processed by the python script. A ROI (region of interest) is extracted from the input frame.

The stencil is the filled version of the ROI frame. In Figure 1d the colors are used as follows:

- Green Left lane arrays.
- Red Mean value of the left lane arrays.
- Blue Right lane arrays.
- Pink Mean value of the right lane arrays.
- Yellow Direction where to go. The lower point is located at the center of the image edge. The upper point is the center between the distant points of both lanes. This is only the case while both lanes are detected. If only one lane is detected the upper point of the yellow line is at the endpoint of the edge moved by a quarter of the picture width. This is estimated to be half of the track width.

This is done by using the Hough line transformation [2]. Normally edge detection is part of the algorithm. In other words the edge image is crucial as input. In our case an edge detection algorithm like Canny is redundant. This is first because of the simplicity of our input frames. Second the ROI frame which is inputted into the Hough line transformation is inverted because the ROI frame is black on white while most edge detection outputs are white on black. By inverting it the resulting input to the transformation is close to edge images. The other input parameters for the python HoughLinesP function are theta, threshold, minLineLength and maxLaneGap.

minLineLength The minimum line length is needed so that only the track edges are detected and processed as lanes. If the value is too low the edge of the right front wheel is detected which can be clearly seen in Figure 1a and 1b.

maxLineGap This describes the distance between the endpoints of both lanes at the picture center and is estimated.

The P function ending indicates that the probabilistic Hough line transform is applied instead of the standard version. It is a more efficient implementation which has a different output, the extremes of the detected lines. Instead of only the function it is helpful to understand how the Hough line transform works.

•TODO: explain

Also a helper function was used to distinguish lines marking left and right lanes. Since the Hough line transformation algorithm walks column-wise through the image from the left top to the right bottom the helper function splits right and left. The left lane is completely detected and finished before the area of the image where the right lane is located is processed by the algorithm.

C. Autonomous Driving

Both subtasks merged together result in autonomous driving. Video 1 shows a short extract from driving.

The video also contains the steering angle values which are also indicated by the yellow direction line as shown in Figure 1d. The autonomous driving script contains the helper function

Video 1: Direction and steering angle while driving.

to calculate the direction which was already explained in Section II-B. If no lane is detected the direction for steering is straight ahead. Another helper function calculates the steering angle from the desired direction which can be used afterwards as the input for the steering command described in Section II-A. To minimize the error rate of the steering angle a helper function stabilizes the angle by comparing it to the previous calculated values. This is useful since the direction differs a lot if only one lane is detected correctly. Video 2 contains some of these errors. While steering left the direction aims a few times to the right lane because as explained before the direction is pointing towards the endpoint of the edge moved by a quarter of the picture width while the left lane is not detected correctly.

Video 2: Dataset with control values.

Video 3 shows the car driving a turn from external view.

Video 3: Driving a turn from external view.

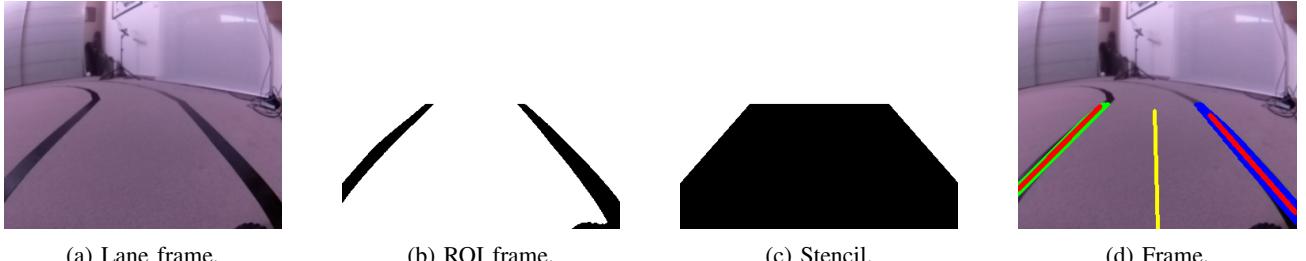


Fig. 1: Lane detection.

•TODO: discuss

D. Benefit

It may be unclear what the benefit from an approach without machine learning might be. Therefore this will now be clarified. The approach is completely detached from donkeycar. Steering and motor are controlled directly through the Arduino PCA controller. Also the code is from scratch in latest versions. This is a great benefit for future work with the car in a teaching environment. Process thinking regarding driving is detached from donkeycar and also machine learning approaches in future work can be done completely without it what marks a difference to our machine learning approaches. Also it clarifies which components are built in the car. This might be unclear since receiving an already built car pushes to box thinking without being aware of the connection between parts or the presence and function of the Arduino controller.

•TODO: (Decorator, Dataclass, “sugaring”)

III. DONKEYCAR AUTOPILOT

Although we decided against using the donkeycar library •TODO: (forgot why), we did experiment with it in an OpenAI gym environment called gym-donkeycar, which simulates a car that a donkeycar application can run on.

Donkeycar 2.5.8 comes with the ability to collect data and train a neural network built-in.

A. Data Collection

At each step the gym environment returns the image that the virtual car sees. The steering angle and throttle value are passed to the gym as a numpy array of length two by the donkeycar application. When recording drive footage, *json* files containing an image file name, steering angle, throttle value and other miscellaneous information are saved with the associated images in a folder. These folders can be used to train a model.

B. Training

Different model types are available under the *donkeycar.parts.keras module*. As the name indicates these model types are built with the Keras API on top of TensorFlow. The available types are:

Categorical takes an image as input and has two categorical outputs for steering and throttle.

The input is passed to a network with the following layers:

- 1) 5x5 Conv-ReLU, 24 filters, stride 2
- 2) 5x5 Conv-ReLU, 32 filters, stride 2
- 3) 5x5 Conv-ReLU, 64 filters, stride 2 or 3x3 Conv-ReLU, 64 filters, stride 1
- 4) 3x3 Conv-ReLU with 64 filters and stride 2 or 3x3 Conv-ReLU, 64 filters, stride 1
- 5) 3x3 Conv-ReLU, 64 filters, stride 1
- 6) Dense-ReLU, 100 outputs
- 7) Dense-ReLU, 50 outputs
- 8) Dense-Softmax, 15 outputs for steering and Dense-Softmax, 20 outputs for throttle

Linear

takes an image as input and has two scalar outputs for steering and throttle. The input is passed to a network with the following layers:

- 1) 5x5 Conv-ReLU, 24 filters, stride 2
- 2) 5x5 Conv-ReLU, 32 filters, stride 2
- 3) 5x5 Conv-ReLU, 64 filters, stride 2
- 4) 3x3 Conv-ReLU, 64 filters, stride 1
- 5) 3x3 Conv-ReLU, 64 filters, stride 1
- 6) Dense-ReLU, 100 outputs
- 7) Dense-ReLU, 50 outputs
- 8) Dense-Linear, 1 output for steering and Dense-Linear, 1 output for throttle

IMU

takes an image and an IMU vector as input and has two scalar outputs for steering and throttle. The image is passed to a Linear network without layers 7 and 8, while the IMU vector is passed to a network with 3 Dense-ReLU layers with 14 outputs each. The outputs of both networks are concatenated and passed to a network with 2 Dense-ReLU layers with 50 outputs each and two separate Dense-Linear layers with 1 output for steering and throttle respectively.

Latent

takes an image as input and has two scalar outputs and an image output. This experimental model type uses convolutional layers to learn a latent vector and transposed convolutional layers to reconstruct an image from the latent vector as well as dense layers to produce the

steering and throttle outputs from the same latent vector.

RNN

takes a sequence of images as input and has one 2D output for steering and throttle. The images are passed to a Linear network whose last Conv-ReLU layer is replaced with a 2x2 MaxPooling layer and the Dense-ReLU layer with 50 outputs is removed entirely. To make use of the sequence of images, each layer is wrapped inside a TimeDistributed layer, which applies the wrapped layers to each input. Two Long Short-Term Memory layers with 128 outputs followed by four Dense-ReLU layers with 128, 64, 10, and 2 outputs respectively complete the network.

3D

takes a sequence of images as input and has one 2D output for steering and throttle. The input is passed to a network with the following layers:

- 1) 3x3x3 3D Conv-ReLU, 16 filters, stride (1,3,3)
- 2) 1x2x2 3D MaxPooling, stride (1,2,2)
- 3) 3x3x3 3D Conv-ReLU, 32 filters, stride (1,1,1)
- 4) 1x2x2 MaxPooling, stride (1,2,2)
- 5) 3x3x3 3D Conv-ReLU, 64 filters, stride (1,1,1)
- 6) 1x2x2 3D MaxPooling, stride (1,2,2)
- 7) 3x3x3 3D Conv-ReLU, 128 filters, stride (1,1,1)
- 8) 1x2x2 3D MaxPooling, stride (1,2,2)
- 9) Dense-BatchNorm-ReLU, 256 outputs
- 10) Dense-BatchNorm-ReLU, 256 outputs
- 11) Dense, 2 outputs

Behaviour

takes an image and a behaviour vector as input and has two categorical outputs. The image is passed to a Linear network without layers 7 and 8, while the Behaviour vector is passed to a network with 3 Dense-ReLU layers with each layers outputs twice its inputs. The outputs of both networks are concatenated and passed to a network with a Dense-ReLU layers with 100 outputs followed by one with 50 outputs and two separate Dense-Linear layers with 15 outputs for steering and 20 outputs for throttle.

Localizer

takes an image as input and has two scalar outputs for steering and throttle and one categorical output for location. The image is passed to a Linear network which has an additional Dense-ReLU output layer for a location output with the number of outputs specified by the user.

IV. DAGGER

A. Description

DAGGER (Dataset Aggregation), developed by Ross et al. [4], is an iterative algorithm, which in each iteration i uses policy π_i to collect in dataset \mathcal{D}_i state-action pairs $(s, \pi^*(s))$, where s is the state visited by the policy π_i and $\pi^*(s)$ is the associated action returned by an expert policy π^* . A classifier $\hat{\pi}_{i+1}$ is then trained on $\mathcal{D} = \mathcal{D}_1 \cup \dots \cup \mathcal{D}_i$ (see Figure 2).

In practice this means an expert, in our case •TODO: (lane follower), takes control of the autonomous car and collects image-action pairs, where the action is a steering angle and a throttle value. A classifier is trained on the available data and used to collect new images, all of which are then labeled by the expert. Those new image-action pairs are combined with the old and used to train the next classifier. The process of data collection, expert labelling, and training is repeated until a satisfactory classifier has been trained.

```

1: Initialize  $\mathcal{D} \leftarrow \emptyset$ 
2: for  $i = 1$  to  $N$  do
3:   Let  $\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}_i$ .
4:   Sample  $T$ -step trajectories using  $\pi_i$ .
5:   Get dataset  $\mathcal{D}_i = \{(s, \pi^*(s))\}$  of visited states by  $\pi_i$ 
       and actions given by expert.
6:   Aggregate datasets:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$ .
7:   Train classifier  $\hat{\pi}_{i+1}$  on  $\mathcal{D}$ .
8: end for
9: return best  $\hat{\pi}_i$  on validation.

```

Fig. 2: DAGGER algorithm

Note, that due to the lack of data and therefore policies in the first iteration, we usually set

$$\beta_i = \begin{cases} 1, & \text{if } i = 1 \\ 0, & \text{else} \end{cases}$$

or

$$\beta_i = \begin{cases} 1, & \text{if } i = 1 \\ p^{i-1}, & \text{else} \end{cases}$$

where p is some probability, to directly utilize the expert as our first policy. Subsequent policies $\hat{\pi}_{i+1}$ are trained on expert labeled states visited by the current policy π_i or a combination of the current policy and the expert policy.

B. Implementation

We used a neural network, implemented in PyTorch, similar to the Linear network used by donkeycar to train the next policy on the expert labeled dataset. For details on the expert implementation see section •TODO: (no ml). We set $\beta_1 = 1$ utilizing the expert in the first iteration and $\beta_{2:n} = 0$ utilizing the trained policy in the remaining iterations.

C. Results

We measure distance travelled by the trained policy until the car veers off track. See table II for results.

TABLE II: Results

iteration	distance travelled
1	x laps

Video 4: Dataset.

Video 5: Model driving in simulation.

V. CONCLUSION

VI. PLACEHOLDER

move to best fitting (sub-)section

REFERENCES

- [1] M. Abadi et al. “Tensorflow: Large-scale machine learning on heterogeneous distributed systems,” arXiv preprint, 2016.
- [2] P. Hough. “Method and means for recognizing complex patterns”, United States Patent Office, 3069654, 1962
- [3] V. Leavers. “Shape Detection in Computer Vision Using Hough Transform”, Springer, London, UK, 1992
- [4] S. Ross et al. “A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning”, Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics, AISTATS 2011, Fort Lauderdale, USA, April 11-13, 2011